

# RISC-V Architecture and Supercomputing

Thomson Kneeland, Tai-Yin Chong, and David  
Berstol

# CISC - Complex Instruction Set Computer

- Predecessor of RISC
- Capable of multistep operations and addressing modes with single instructions
- Complete task with minimal code and rely on hardware
  - Minimizes compiler work with less assembly code
  - Build task directly into hardware
- Multiclock complex instructions
- High cycles per second
- Memory to memory “load” and “store” incorporated in instructions
- Intel x86 only chip that retains CISC architecture

# RISC- Reduced Instruction Set Computer

- Developed in late 70's: IBM, Stanford, UC Berkeley
- RISC ISA allows the processor to average nearly 1 instruction per cycle
  - Instruction **S**et **A**rchitecture is the interface between software and hardware
- ALU operations between registers only (load/store architecture)
  - In CISC arithmetic, one operand could be in memory and another in a register
  - Large number of registers to minimize interaction with memory
- Pipelining - allows for simultaneous execution of parts
- OS and app programmers can develop code easily
- MIPS is a reduced instruction set developed at Stanford in 1984
  - Lowers compiler to hardware level

# RISC V CPU Architecture

- Entered the market in 2014
- Originally designed to support education and computer architecture research
- Allows for efficient implementation by avoiding “overarchitecturing”
- Open source and available to public
- Potentially faster and more efficient than x86 and ARM, given good implementation
  - RISC-V Rocket core twice as energy efficient as ARM counterpart Cortex-A5
- Focuses on performance and power efficiency

# Implementations of RISC V ISA

- Cloud computers
- Mobile phones
  - Samsung
- Embedded systems
- Supercomputers
- Storage devices
  - Western Digital
- GPUs
  - Nvidia



# Migration to RISC V

- Lack of alternative RISC ISA
  - ARM bought out by Softbank, MIPS difficult to use due to intellectual property issues
  - Low competitive RISC ISA to ARM and MIPS until RISC V
- The RISC V foundation allows companies to manage the instruction set as a team
  - Berkeley Architecture Research, Google, Samsung, NVidia, and many more
- Western Digital
  - Recently announced their processors transitioning over to RISC V

①

Base Integer Instructions: RV32I, RV64I, and RV128I					RV Privileged Instructions						
Category	Name	Fmt	RV32I Base		+RV{64,128}		Category	Name	RV mnemonic		
Loads	Load Byte	I	LB	rd,rs1,imm			CSR Access	Atomic R/W	CSRrw rd,csr,rs1		
	Load Halfword	I	LH	rd,rs1,imm				Atomic Read & Set Bit	CSRrs rd,csr,rs1		
	Load Word	I	LW	rd,rs1,imm	L{D Q}	rd,rs1,imm		Atomic Read & Clear Bit	CSRrc rd,csr,rs1		
	Load Byte Unsigned	I	LBU	rd,rs1,imm				Atomic R/W Imm	CSRrwi rd,csr,imm		
	Load Half Unsigned	I	LHU	rd,rs1,imm	L{W D}U	rd,rs1,imm		Atomic Read & Set Bit Imm	CSRrsi rd,csr,imm		
Stores	Store Byte	S	SB	rs1,rs2,imm			Atomic Read & Clear Bit Imm	CSRrci rd,csr,imm			
	Store Halfword	S	SH	rs1,rs2,imm			Change Level	Env. Call	ECALL		
	Store Word	S	SW	rs1,rs2,imm	S{D Q}	rs1,rs2,imm		Environment Breakpoint	EBREAK		
Shifts	Shift Left	R	SLL	rd,rs1,rs2	SLL{W D}	rd,rs1,rs2		Environment Return	ERET		
	Shift Left Immediate	I	SLLI	rd,rs1,shamt	SLLI{W D}	rd,rs1,shamt	Trap Redirect to Supervisor	MRTS			
	Shift Right	R	SRL	rd,rs1,rs2	SRL{W D}	rd,rs1,rs2		Redirect Trap to Hypervisor	MRTS		
	Shift Right Immediate	I	SRLI	rd,rs1,shamt	SRLI{W D}	rd,rs1,shamt	Hypervisor Trap to Supervisor	MRTS			
	Shift Right Arithmetic	R	SRA	rd,rs1,rs2	SRA{W D}	rd,rs1,rs2	Interrupt	Wait for Interrupt	WFI		
	Shift Right Arith Imm	I	SRAI	rd,rs1,shamt	SRAI{W D}	rd,rs1,shamt		MMU	Supervisor FENCE	SFENCE.VM rs1	
Arithmetic	ADD	R	ADD	rd,rs1,rs2	ADD{W D}	rd,rs1,rs2	Optional Compressed (16-bit) Instruction Extension: RVC				
	ADD Immediate	I	ADDI	rd,rs1,imm	ADDI{W D}	rd,rs1,imm					
	SUBtract	R	SUB	rd,rs1,rs2	SUB{W D}	rd,rs1,rs2	Category	Name	Fmt		
	Load Upper Imm	U	LUI	rd,imm			RVC		RVI equivalent		
Logical	Add Upper Imm to PC	U	AUIPC	rd,imm			Loads	Load Word	CL	C.LW rd',rs1',imm	LW rd',rs1',imm*4
	XOR	R	XOR	rd,rs1,rs2				Load Word SP	CI	C.LWSP rd,imm	LW rd,sp,imm*4
								Load Double	CL	C.LD rd',rs1',imm	LD rd',rs1',imm*8
	OR Immediate	R	ORI	rd,rs1,imm				Load Double SP	CI	C.LDSP rd,imm	LD rd,sp,imm*8
	AND	R	AND	rd,rs1,rs2				Load Quad	CL	C.LQ rd',rs1',imm	LQ rd',rs1',imm*16
	AND Immediate	I	ANDI	rd,rs1,imm				Load Quad SP	CI	C.LQSP rd,imm	LQ rd,sp,imm*16

# Sample RISC V program

Editor

Simulator

```
1 ## program that finds value of (2+3)*4, enters in register t4|
2 li t0, 2
3 li t1, 3
4 li t2, 4
5 add t3, t0, t1
6 mul t4, t2, t3
7
```



# Output to register t4

Editor

Simulator

Run

Step

Prev

Reset

Dump

Machine Code	Basic Code	Original Code
0x00200293	addi x5 x0 2	li t0, 2
0x00300313	addi x6 x0 3	li t1, 3
0x00400393	addi x7 x0 4	li t2, 4
0x00628e33	add x28 x5 x6	add t3, t0, t1
0x03c38eb3	mul x29 x7 x28	mul t4, t2, t3

s5  
(x21)

0

s6  
(x22)

0

s7  
(x23)

0

s8  
(x24)

0

s9  
(x25)

0

s10  
(x26)

0

s11  
(x27)

0

t3  
(x28)

5

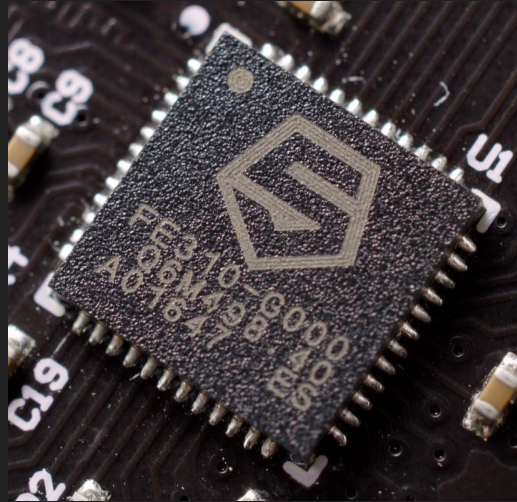
t4  
(x29)

20

t5

0

# The Processors

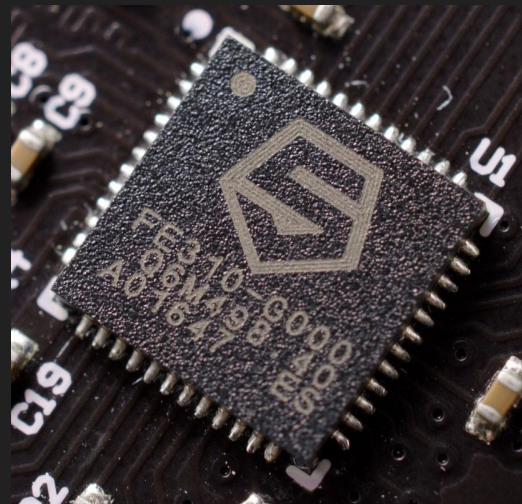
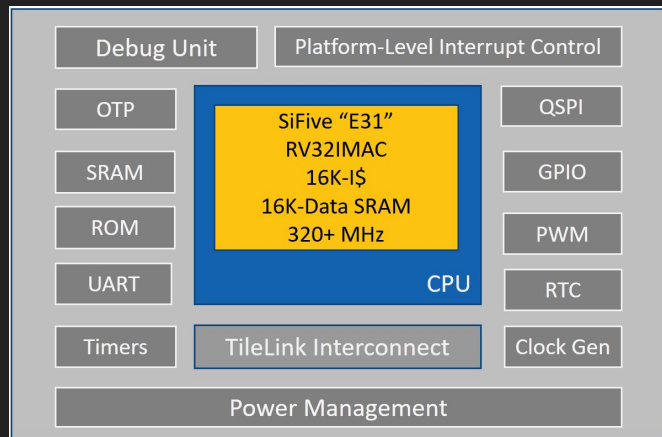


# Esperanto Technologies

- Announced future plans to develop energy-efficient computing solutions for AI and Machine Learning applications
  - Based on RISC-V ISA
- Advantage of using RISC-V
  - Over four thousand full 64-bit cores on a single chip
  - Growing software base of operating systems, compilers and applications
  - Ability to deliver world class TeraFlop levels of computing
  - Increased software availability

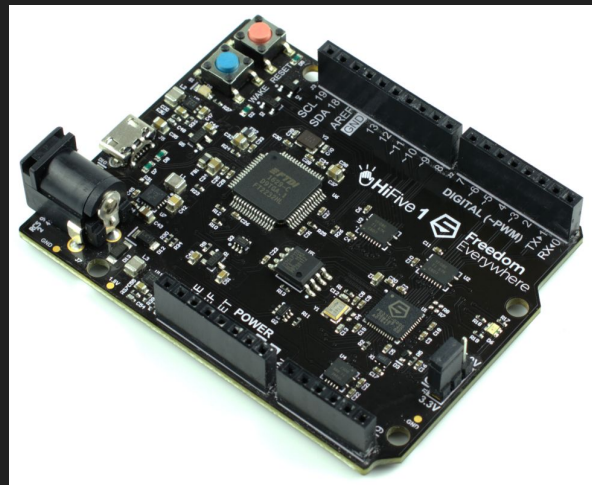
# Freedom E310

- Industry's first commercially available RISC-V SoC
- Released November 29th, 2016
- Developed by SiFive, startup in Silicon Valley
  - First company to produce a chip implementing RISC-V ISA



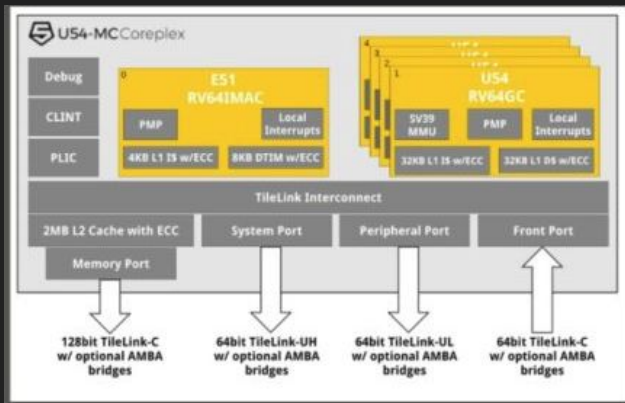
# HiFive1

- Arduino-Compatible development kit
- RISC-V based
  - Features the Freedom E310
- Available for \$59
- Will allow a variety of designers and system architects to create their own products



# U54-MC Coreplex IP

- World's first 64 bit quad-core RISC-V open source processor in Oct 17
- Created by SiFive
- Brings RISC-V into Linux processing
- Four U54 CPUs and a single E51 CPU; running at 1.5GHz each
- 100 prototypes available at \$100,000



U54-MC Coreplex

ARM Cortex-A35

RV64GC M + S + U Mode	ARMv8-A, AArch32, AArch64
16 bit instructions	AArch32 only
Physical Memory Protection (PMP) and MMU	None, MMU only
Real-time capable	Not applicable
E51 Monitor Core	Requires additional IP
Integrated interrupt controller	Requires additional IP

Image: SiFive

# Nvidia Falcon

- Nvidia GPUS have featured a proprietary microcontroller called “Falcon”
  - Fast Logic Controller
  - Protects hardware from misprogramming/malicious software
  - 32 bit address range
  - No caches until recent lcache
- Next generation microcontroller is being built on RISC-V
  - More capabilities include greater address width and configurable cache
  - Security support, software toolchains and more
  - Backward compatibility opens up to 3rd party programming
  - Can mix and match internally/externally developed cores
- NVidia will use RISC-V in many of its future products

# NVidia compared pros and cons

Item	Requirement	ARM A53	ARM A9	ARM R5	RISC-V Rocket	NV RISC-V
Core perf	>2x falcon	Yes	Yes	Yes	Yes	Yes
Area (16ff)	<0.1mm <sup>2</sup>	No	No	Yes	Yes	Yes
Security	Yes	TZ	TZ	No	Yes	Yes
TCM	Yes	Yes	No	Yes	No	Yes
L1 I/D \$	Yes	Yes	Yes	Yes	Yes	Yes
Addressing	64bit	Yes	No	No	Yes	Yes
Extensible ISA	Yes	No	No	No	Yes	Yes
Safety (ECC/Parity)	Yes	Yes	Yes	Yes	Yes	Yes
Functional Simulation model	Yes	Yes	No	No	No	Yes



RISC-V surpassed all in criteria



# Supercomputers



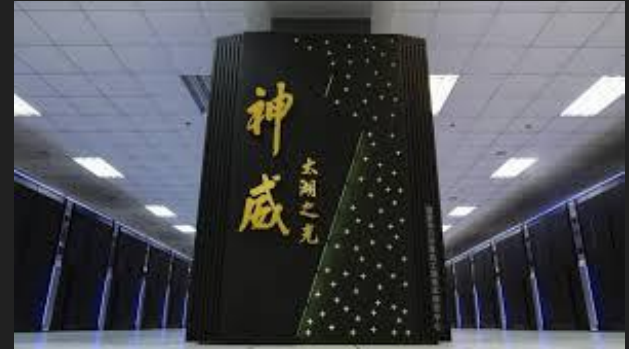
Tsubame 3.01, Japan

# Supercomputers and processor architecture

- Massively Parallel Processing (MPP)
  - Each processor has its own memory and operating system
  - All connected and run in parallel for unified system
  - Can all work on same task
- Early models designed by Cray Research in 1960's
- Computer cluster (local location)
- Distributed computing (BOINC, SETI, etc)
- Can utilize CPUs or a combination of CPUs and GPUS
- Boundaries include power consumption, cooling needs, memory address problems

# Top Supercomputers 11/13/2017

	NAME	COUNTRY	R <sub>MAX</sub> PFLOP/S	POWER MW
1	Sunway TaihuLight	China	93.0	15.4
2	Tianhe-2 (Milkyway-2)	China	33.9	17.8
3	Piz Daint	Switzerland	19.6	2.27
4	Gyokou	Japan	19.1	1.35
5	Titan	USA	17.6	8.2

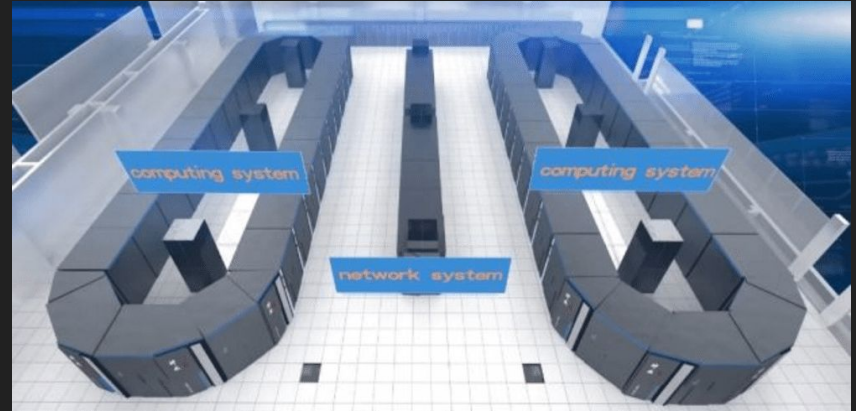


Sunway Taihu Light

# #1 Sunway TaihuLight

## Massively Parallel Processor Array

Distributed memory accessed locally;  
each processor strictly encapsulated



- Benchmark rating of 93 peta flops
- 40,960 SW26010 64-bit RISC Sunway processors
- Each chip has 256 cores, for a total of 10,649,600 CPU cores in system
- Runs on its own operating system based on Linux
- “Network on a chip” vs traditional cache hierarchy

National Research Center in Wuxi China

- Chinese built 64-bit RISC processor, mystery architecture

## #2 Tianhe-2 (TH-2)

- Previous fastest supercomputer, also Sunway, completed at end of 2013
  - 34 petaflop speed in Guangzhou
  - Team of 1,300 scientists and engineers to bring to fruition
  - 3,120,000 cores
  - Chinese built 64-bit RISC processor, mystery architecture
  - CPU processor memory = 1.34 PiB Pebibytes
  - State sponsored
  - 33,860 trillion calculations/second
  - Power consumption: 17,800 kW
  - Cost over \$385 million
  - Occupies 7,750 square ft



# #5 Titan

- Oak Ridge Laboratories in Oak Ridge, Tennessee
  - Sponsored by DOE and NOAA
  - Built by American supercomputer manufacturer Cray
  - Benchmark Rating of 17.59 PetaFLOPS
  - Was #1 in 2012 until overtaken by Tianhe-2
  - Includes GPUs and CPUs
    - 18,688 AMD Opteron 16 core CPUs
    - 18,688 Nvidia Tesla GPUs
  - \$97 million to build
  - Scientific computing, climate models  
Cryptographic analysis



# The future: IBM “Summit”

- In development at Oak Ridge and expected to be operational in 2018
- Uses IBM Power9 CPU and NVIDIA Volta GPU
  - Volta: 100 Teraflops/s deep learning
  - 21 billion transistors
  - CUDA and Tensor Cores: optimized for Deep Learning/AI
  - Should top Titan’s performance by 5-10X
  - >40 TeraFlops
  - Total system memory >10 PebiBytes
  - Will only require 4,600 processors vs Titan’s 18,688





# IBM “Sierra”

- Being built for Livermore for assessment of nuclear weapon science/calculations
- Also IBM built
- Uses IBM Power 9 processors and NVidia’s Volta Graphics CPU
  - Will operate at 5-7X the present Sequoia supercomputer installed
  - 125 PetaFlop/s peak
  - 5 times more power efficient





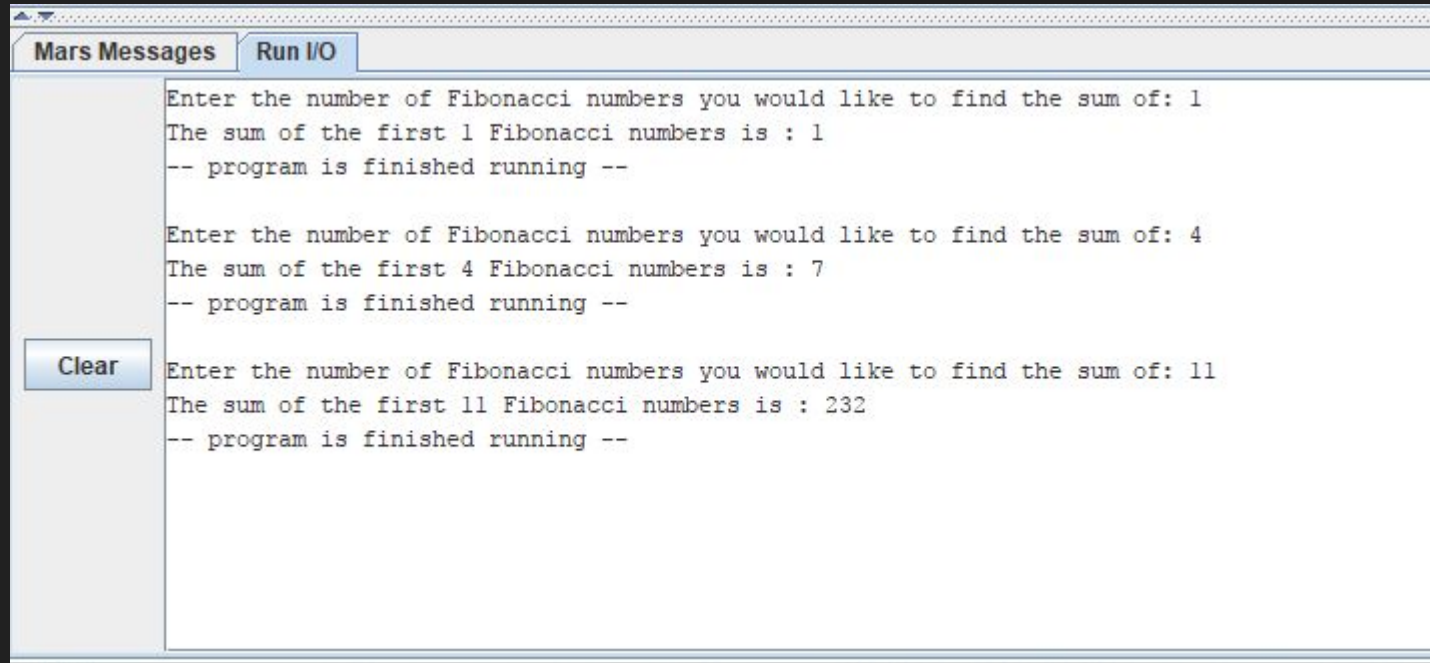
# Summary

- Overall, the industry is moving forward in terms of energy efficiency, size, and computing power
- Parallel Processing is the order of the day for complex calculations and modelling
- RISC and CISC technologies have actually been converging
- RISC-V could usher in increased efficiency for the future; especially as it is available for public utilization
- RISC-V and open source paves the way for increased collaboration for research without being limited by patents and corporate control/licensing

# Appendix: Loop: Sum of Fibonacci Numbers (MIPS)

```
32      la $a0, msg # output message to console
33      li $v0, 4
34      syscall
35
36      li $v0, 5 #Read user input integer
37      syscall
38
39      move $t6, $v0 # store user input n
40      sub $t5,$t6, 1 #subtract one for iteration
41
42 loop:
43     bgt $t3, $t5, exit # loop up to user input number
44     add $t2, $t0, $t1 # add previous two terms and store
45     add $t4, $t4, $t2 # add new term to final sum
46     move $t0, $t1 # store two terms for next iteration
47     move $t1, $t2
48     addi $t3, $t3, 1 # increment counter
49     j loop
50
51 exit:
52     la $a0, msg1 # output message to console
53     li $v0, 4
54     syscall
55
56     move $a0, $t6 # output user input n to console
57     li $v0, 1
58     syscall
59
60     la $a0, msg2 # output second part of message
61     li $v0, 4
62     syscall
63
64     move $a0, $t4 # print final sum to console
65     li $v0, 1
66     syscall
67
```

# Output



The screenshot shows a window titled "Mars Messages" with a "Run I/O" button. The window contains three sets of text, each representing a program execution. A "Clear" button is located on the left side of the window.

```
Enter the number of Fibonacci numbers you would like to find the sum of: 1
The sum of the first 1 Fibonacci numbers is : 1
-- program is finished running --

Enter the number of Fibonacci numbers you would like to find the sum of: 4
The sum of the first 4 Fibonacci numbers is : 7
-- program is finished running --

Enter the number of Fibonacci numbers you would like to find the sum of: 11
The sum of the first 11 Fibonacci numbers is : 232
-- program is finished running --
```

# Register to Memory (MIPS)

```
1 #Tai-Yin Chong
2 #CSIT 230
3 #PROJECT
4 #PROBLEM 3
5 .data
6 Num0: .word 3
7 Num1: .word 4
8 msg0: .asciiz "The sum of the two integers is: "
9 .text
10 .globl main
11 main:
12     la $s0, Num0
13     la $s1, Num1
14
15     la $a0, msg0
16     li $v0, 4
17     syscall
18
19     lw $t0, 0($s0)
20     lw $t1, 0($s1)
21     add $a0, $t0, $t1
22     li $v0, 1
23     syscall
24
25     li $v0, 10
26     syscall
27
28
29
```